

**Exercise 6.1.** Give three examples from the RISC-V architecture of each of the following design principles: (1) regularity supports simplicity; (2) make the common case fast; (3) smaller is faster; and (4) good design demands good compromises. Explain how each of your examples exhibits the design principle.

**Solution:**

- (1) In Section 6.2.1, the `add` and `sub` instructions were introduced, and they both have the same format: `mnemonic source1 source1 destination`. The format is consistent with nearly all RISC-V binary operations, making them predictable and their binary encoding consistent and simple. For example, the `func7` bits are all 0 for all binary operations with an `op` code of 51; they differ only in their `funct3` control bits. Another example is that operands all 32 bits in RISC-V 32, in contrast with x86-32 where operands can be 8-bit, 16-bit, or 32-bit. Also instructions in x86-64 vary in size, whereas in RISC-V they are uniformly 16 bits, as mentioned in Section 6.8.5.
- (2) As discussed in 6.2.1, RISC-V has a relatively small set of simple instructions that are commonly used, in contrast with CISC architectures that provide complex operations and therefore add overhead to simple instructions. In Section 6.8.6, the author mentions that x86 contains string operations that are usually slower than performing the equivalent operation with a series of simple instructions. By supporting simple byte operations that are fast, RISC-V can reap these benefits.
- (3) RISC-V32 has 32 registers, which are sufficient to perform several operations without having to add local variables to the stack in the case of small functions. Since the stack relies on larger memory, accessing is slower, so avoiding this is best.
- (4) In Section 6.4, the author mentions compromise regarding the length of the instruction encoding; some instructions do not require 32 bits, but they are encoded as such anyway to support simplicity. Moreover, instead of a single instruction format, there are four to provide enough flexibility. Another example is the bit-swizzling of the immediate encodings discussed in 6.4.5. Though complicated, it ensures that other instruction fields are consistent and simple.

**Exercise 6.2.** The RISC-V architecture has a register set that consists of 32 32-bit registers. Is it possible to design a computer architecture without a register set? If so, briefly describe the architecture, including the instruction set. What are the advantages and disadvantages of this architecture over the RISC-V architecture?

**Solution:** According to Section 6.2.2, registers exist to support access to operands quickly so that they can run fast; they do this by providing a small address space. Therefore, it is possible to not use a register set and use memory. One way would be to reserve the first 32

address locations of memory for the same functions as the corresponding registers. According to Section 5.5.4, memory is commonly implemented with DRAM, whereas a register set is implemented with SRAM. DRAM latency is longer than SRAM because its bitline is not actively driven by a transistor; it fundamentally has lower throughput because data must be refreshed periodically and after a read. Moreover, latency and throughput depend on memory size, and therefore since DRAM is typically very large in comparison to a small register set, this means relying on just memory would be too slow. One advantage is that it generally takes less transistors to build DRAM in comparison to SRAM, making them cheaper and less power hungry.

**Exercise 6.3.** Write the following strings using ASCII encoding. Write your final answers in hexadecimal.

- (a) hello there
- (b) bag o'chips
- (c) To the rescue!

**Solution:**

- (a) Referring to an ASCII table encoding, we see that the encoding is:

HEX	68	65	6c	6c	6f	20	74	68	65	72	65
String	h	e	l	l	o		t	h	e	r	e

- (b)

HEX	62	61	67	20	27	6f	63	68	69	70	73
String	b	a	g		o	'	c	h	i	p	s

- (c)

HEX	54	6f	20	74	68	65	20	72	65	73	63	75	65	21
String	T	o		t	h	e		r	e	s	c	u	e	!

**Exercise 6.4.** Repeat Exercise 6.3 for the following strings.

1. Cool
2. RISC-V
3. boo!

**Solution:**

- (a)

HEX	43	6f	6f	6c
String	C	o	o	l

(b)

HEX	52	49	53	43	2d	56
String	R	I	S	C	-	V

(c)

HEX	62	6f	6f	21
String	b	o	o	!

**Exercise 6.5.** Show how the strings in Exercise 6.3 are stored in byte-addressable memory starting at memory address 0x004F05BC. The first character of the string is stored at the lowest byte address (in this case, 0x004F05BC). Clearly indicate the memory address of each byte.

**Solution:** I am using word addresses that increase upwards. The last character is therefore at the top.

(a)

hello there	
Byte	Word Address
65	0x004F05C6
72	0x004F05C5
65	0x004F05C4
68	0x004F05C3
74	0x004F05C2
20	0x004F05C1
6f	0x004F05C0
6c	0x004F05BF
6c	0x004F05BE
65	0x004F05BD
68	0x004F05BC

(b)

bag o'chips		
Byte	Word	Address
73		0x004F05C6
70		0x004F05C5
69		0x004F05C4
68		0x004F05C3
63		0x004F05C2
6f		0x004F05C1
27		0x004F05C0
20		0x004F05BF
67		0x004F05BE
61		0x004F05BD
62		0x004F05BC

(c)

To the rescue!		
Byte	Word	Address
21		0x004F05C9
65		0x004F05C8
75		0x004F05C7
63		0x004F05CD
73		0x004F05C5
65		0x004F05C4
72		0x004F05C3
20		0x004F05C2
65		0x004F05C1
68		0x004F05C0
74		0x004F05BF
20		0x004F05BE
6f		0x004F05BD
54		0x004F05BC

**Exercise 6.5.** Repeat Exercise 6.5 for the strings in Exercise 6.4.

**Solution:** As in Exercise 6.5, I am using word addresses that increase upwards. The last character is therefore at the top.

(a)

Cool		
Byte	Word	Address
6c		0x004F05BF
6f		0x004F05BE
6f		0x004F05BD
43		0x004F05BC

(b)

RISC-V		
Byte	Word	Address
56	0x004F05C1	
2d	0x004F05C0	
43	0x004F05BF	
53	0x004F05BE	
49	0x004F05BD	
52	0x004F05BC	

(c)

boo!		
Byte	Word	Address
21	0x004F05C3	
6f	0x004F05C2	
6f	0x004F05C1	
62	0x004F05C0	

**Exercise 6.7.** The `nor` instruction is not a part of the RISC-V instruction set because the same functionality can be implemented using existing instructions. Write a short assembly code snippet that has the same functionality: `s3 = s4 NOR s5`. Use as few instructions as possible.

**Solution:** Recall that NOR is short for NOT OR, so  $A \text{ NOR } B$  means  $\overline{A + B}$ , where  $+$  is the OR operator and the overline is the NOT operator. The truth table is below:

A	B	$\overline{A + B}$
0	0	1
0	1	0
1	0	0
1	1	0

Recalling that NOT can be implemented with `xori` with `-1` as the immediate, do precisely an OR followed by a NOT:

---

```
or s6, s4, s5
xori s3, s6, -1
```

---

**Exercise 6.8.** The `nand` instruction is not a part of the RISC-V instruction set because the same functionality can be implemented using existing instructions. Write a short assembly code snippet that has the same functionality: `s3 = s4 NAND s5`. Use as few instructions as possible.

**Solution:** Recall that NAND is short for NOT AND, so  $A \text{ NAND } B$  means  $\overline{A \cdot B}$ , where  $\cdot$  is the AND operator and the overline is the NOT operator. The truth table is below:

A	B	$\overline{A \cdot B}$
0	0	1
0	1	1
1	0	1
1	1	0

Recalling that NOT can be implemented with `xori` with `-1` as the immediate, do precisely an AND followed by a NOT:

---

```
and s6, s4, s5
xori s3, s6, -1
```

---

**Exercise 6.9.** Convert the following high-level code snippets into RISC-V assembly language. Assume that the (signed) integer variables `g` and `h` are in register `a0` and `a1`, respectively. Clearly comment your code:

(a)

---

```
if (g > h)
    g = g + 1;
else
    h = h - 1;
```

---

(b)

---

```
if (g <= h)
    g = 0;
else
    h = 0;
```

---

**Solution:**

- (a) We can use a conditional jump, `bge`, to go to the `else` branch when `h >= g`. To ensure we do not execute the `else` branch if we take the `if` branch, we use an unconditional jump with the `j` instruction to the label after the last instruction in the `else` branch. To add the constant, we can use the `addi` command. The assembly follows:

---

```
bge    a1, a0      # if h >= g goto else at .L1
addi   a0, a0, 1   # g = g + 1
j      .L2         # Skip else branch
.L1:
addi   a1, a1, -1, # h = h - 1
.L2:
```

---

- (b) We can use `blt` to go to the `else` branch if `h > g` (equivalently, `g > h`).

---

```

    blt    a1, a0      # if h < h goto else at .L1
    addi   a0, zero, 0 # g = 0
    j      .L2         # Skip else branch
.L1:
    addi   a1, zero, 0 # h = 0
.L2:

```

---

**Exercise 6.10.** Repeat Exercise 6.9 for the following code snippets:

(a)

---

```

if (g >= h)
    g = g + h;
else
    g = g - h;

```

---

(b)

---

```

if (g < h)
    h = h + 1;
else
    h = h * 2;

```

---

**Solution:**

(a)

---

```

    blt    a1, a0      # if h < g goto else at .L1
    add    a0, a0, a1  # g = g + h
    j      .L2         # Skip else branch
.L1:
    sub    a0, a0, a1  # g = g - h
.L2:

```

---

(b) Instead of multiplying by 2 with the `mul` instruction, I used `sll` to shift left by 1 bit, which is equivalent.

---

```

    bge    a0, a1      # if g >= h goto else at .L1
    addi   a1, a1, 1    # h = h + 1
    j      .L2         # Skip else branch
.L1:
    sll    a1, a1, 1    # h = h * 2
.L2:

```

---

**Exercise 6.11.** Convert the following high-level code snippet into RISC-V assembly. Assume that the base address of `array1` and `array2` are held in `t1` and `t2` and that `array2` array is initialized before it is used. Use as few instructions as possible. Clearly comment your code.

---

```
int i;
int array1[100];
int array2[100];
# ... more code
for (i = 0; i < 100; i = i + 1)
    array1[i] = array2[i];
```

---

**Solution:**

---

```
    addi    sp,sp,-16 # Allocate space on stack
    sw      s0,12(sp) # Store saved register on stack
    sw      s1,8(sp)  # Store saved register on stack
    addi    s0,t1,0    # Move array1 to saved register.
    addi    s1,t2,0    # Move array 2 to saved register.
    addi    t0,zero,0  # i = 0
    addi    t1,zero,100 # t1 = 100
.LOOP:
    bge     s0,t0,.AFTER_LOOP
    slli    t2,t0,2     # t2 = i * 4
    add     t3,s1,t2     # Address of array2[i]
    lw      t4, 0(t3)    # value stored at array2[i]
    addi    t5,s0,t2     # Address of array1[i]
    sw      t4, 0(t5)    # Store value of array2[i] into array1[i]
    addi    t0,t0,1     # i = i + 1
    j       .LOOP       # Repeat
.AFTER_LOOP:
    lw      s0,12(sp)    # Restore saved registers
    lw      s1,8(sp)
    addi    sp,sp,16     # Deallocate space on stack
```

---

**Exercise 6.12.** Repeat Exercise 6.11 for the following high-level code snippet. Assume that the `temp` array is initialized before it is used and that `t3` holds the base address of `temp`.

---

```
int i;
int temp[100];
...
for (i = 0; i < 100; i = i + 1)
    temp[i] = temp[i] * 128;
```

---

**Solution:**



---

```

    addi    sp,sp,-16 # Allocate space on stack
    sw      s0,12(sp) # Store saved register on stack
    addi    s0,t3,0   # Move temp to saved register.
    addi    t0,zero,0 # i = 0
    addi    t1,zero,100 # t1 = 100
.LOOP:
    bge     s0,t0,.AFTER_LOOP
    slli    t2,t0,2    # t2 = i * 4
    add     t2,s0,t2    # Address of temp[i]
    lw      t3, 0(t2)  # Value stored at temp[i]
    slli    t3,t3,7    # temp[i] * 128
    sw      t3, 0(t2)  # temp[i] * 128 at temp[i]
    addi    t0,t0,1    # i = i + 1
    j       .LOOP      # Repeat
.AFTER_LOOP:
    lw      s0,12(sp)  # Restore saved registers
    addi    sp,sp,16   # Deallocate space on stack

```

---

**Exercise 6.13.** Write RISC-V assembly code for placing the following immediate (constants) in `s7`. Use a minimum number of instructions.

- (a) 29
- (b) -214
- (c) -2999
- (d) 0xABCDE000
- (e) 0xEDCBA123
- (f) 0xEEEEEFAB

**Solution:**

- (a) `addi s7, zero, 29`
- (b) `addi s7, zero, -214`
- (c) -2999 is 1111010001001001 as a 16-bit two's complement number. The upper 4 bits are 1111, and the lower 4 bits are 010001001001. Since the sign extension of the lower 12 bits would not add 1s, we need to use `lui` to add them:

---

```

    lui s7,-1          # adds all 1s to upper 20 bits
    addi s7,0b010001001001 # Append lower 12 bits

```

---

- (d) `lui s7, 0xABCDE`

- (e) The lower 12 bits represented by 0x123 are not signed, we can use `lui` without worrying about the sign extension due to `addi`

---

```
lui s7, 0xEDCBA    # Upper 12 bits
addi s7,s7,0x123    # Sign-extension will subtract 1 from upper bits
```

---

- (f) Since the lower 12 bits 0xFAB is negative, the sign extension will add 1s to the higher 20 bits, so we must add an extra 1 to the upper immediate

---

```
lui s7,0xEEEEF000 # Store 1 bit higher than 0xEEEEE
addi s7,s7,0xFAB   # Sign-extension will subtract 1 from upper bits.
```

---

**Exercise 6.14.** Repeat Exercise 6.13 for the following immediates.

- (a) 47
- (b) -349
- (c) 5328
- (d) 0xBBCCD000
- (e) 0xFEEBC789
- (f) 0xCCAAB9AB

**Solution:**

- (a) `addi s7,zero,47`
- (b) `addi s7,zero,-349`
- (c) Since 5328 is 01010011010000 in 14-bit two's complement binary representation, we need to load the upper bits 01 first, adding 1 to it in anticipation of the sign extension from `addi`:

---

```
lui s7,0b10          # 1 higher than necessary
addi s7,s7,0b010011010000 # Lower 12 bits
```

---

- (d) `lui s7,0xBBCCD`
- (e) Since 0x789 is positive, the sign extension will not add 1s to the higher 20 bits, so we do not add 1 to the upper immediate.

---

```
lui s7,0xFEEBC      # 1 higher than necessary
addi s7,s7,0x789    # sign extension subtracts 1
```

---

- (f) This time we do need to add 1 to the upper immediate since the leading bit of the lower immediate 0x9AB is 1.

---

```
lui s7,0xCCAAC    # 1 higher than necessary
addi s7,s7,9AB    # Sign extension subtracts 1e
```

---

**Exercise 6.15.** Write a function in a high-level language for

---

```
int find42(int array[], int size)
```

---

`size` specifies the number of elements in `array`, and `array[]` specifies the base address of the array. The function should return the index number of the first entry that holds the value 42. If no array entry is 42, it should return `-1`. Clearly comment your code.

**Solution:**

---

```
int find42(int array[], int size)
{
    int i;
    for (i = 0; i < size; i++)
        if (array[i] == 42)
            return i;          // Return index matching 42
    return -1;                 // Not found
}
```

---

**Exercise 6.16.** The high-level `strcpy` (string copy) copies the character string `src` to the character string `dst`

---

```
// C code
// WARNING: This is vulnerable to a buffer overflow attack.
void strcpy(char dst[], char src[]) {
    int i = 0;
    do {
        dst[i] = src[i];
    } while (src[i++]);
}
```

---

- (a) Implement the `strcpy` function in RISC-V assembly. Use `t0` for `i`.
- (b) Draw a picture of the stack before, during, and after the `strcpy` function call. Assume `sp = 0xFFC000` just before `strcpy` was called.

**Solution:**

- (a)

---

```

    # void strcpy(char dst[], char src[]);
strcpy:
    addi    t0,zero,0        # i = 0
.loop:
    slli    t1,t0,2          # t1 = i * 4
    addi    t2,a0,t1         # address of dst[i]
    addi    t3,a1,t1         # address of src[i]
    lw      t3, 0(t3)        # value of src[i]
    sw      t3, 0(t2)        # store src[i] into dst[i]
    addi    t0,t0,1          # i = i + 1
    bne     t3, .loop        # if src[i] != 0 goto .loop
.return
    jr      ra               # return

```

---

- (b) My implementation used only temporary registers so the stack frame for `strcpy` is empty. The stack frame for the calling function contains the elements of `src` and `dst`, since local arrays are stored on the stack.

**Exercise 6.17.** Convert the high-level function from Exercise 6.15 into RISC-V assembly code. Clearly comment your code.

**Solution:**

---

```

    # int find42(int array[], int size)
find42:
    addi    t0, zero, 0      # i = 0
    addi    t1, zero, 42     # t1 = 42
.loop:
    bge     t0,a1,.not_found # if i >= size goto .not_found
    slli    t2,t0,2          # t2 = i * 4
    add     t2,a0,t2         # add base address of array to t2
    lw      t3, 0(t2)        # get value at array[i]
    be      t3,t1            # if array[i] == 42 goto .found
    addi    t0,t0, 1         # i = i + 1
    j       .loop            # Repeat
.found:
    add     a0,zero,t0       # found, return i
    jr      ra               # return
.not_found:
    addi    a0, zero,-1      # not found, returning -1
    jr      ra               # return

```

---

**Exercise 6.18.** Consider the RISC-V assembly code below. `func1`, `func2`, and `func3` are nonleaf functions. `func4` is a leaf function. The code is not shown for each function, but the comments indicate which registers are used within each function. You may assume that the functions do not need to save any nonpreserved registers on their stacks.

---

```

0x00091000 func1: ... # func1 uses t2-t3, s4-s10
...
0x00091020      jal func2
...
0x00091100 func2: ... # func2 uses a0-a2, s0-s5
...
0x0009117C      jal func3
...
0x00091400 func3: ... # func3 uses t3, s7-s9
...
0x00091704      jal func4
...
0x00093008 func4: ... # func4 uses s10-s12
...
0x00093118      jr ra
...

```

---

- (a) How many words are the stack frames of each function?
- (b) Sketch the stack after `func4` is called. Clearly indicate which registers are stored where on the stack and mark each of the stack frames. Give values where possible. Assume that `sp = 0xABC124` just before `func1` is called.

**Solution:**

- (a) By the callee-saved rule, `func1` pushes `s4` through `s10` and `ra` on the stack so it can use those registers. By the instructions, functions do not need to save non-preserved registers, so `t2-t3` are not saved. Therefore, `func1` has a stack frame that is 8 words deep.

By the callee-saved rule, `func2` pushes `s4`, `s5`, and `ra` on the stack, since they are callee-preserved registers. Its stack frame is three words deep.

`func3` pushes `s7-s9` and `ra` on the stack, making its stack frame 4 words deep.

`func4` pushes `s10` on the stack, making its stack frame 1 word deep. Note that the return address register need not be pushed because it is a leaf function.

- (b) See the Figure 1:

**Exercise 6.19.** Each number in the Fibonacci series is the sum of the two previous two numbers. The following table lists the first few numbers in the sequence, `fib(n)`.

$n$	1	2	3	4	5	6	7	8	9	10	11
<code>fib(n)</code>	1	1	2	3	5	8	13	21	34	55	89

- (a) What is `fib(n)` for  $n = 0$  and  $n = -1$ ?




Figure 1: Exercise 6.18: Stack frame after calling `func4`

- Write a function called `fib` in a high-level language that returns the Fibonacci number for any non-negative value of `n`. Hint: You probably will want to use a loop. Clearly comment your code.
- Convert the high-level function of part (b) into RISC-V assembly code. Add comments after every line that explain clearly what it does. Use a simulator to test your code on `fib(9)`. (See the preface for links to a RISC-V simulator).

### Solution:

- `fib(n)` is not defined for  $n = 0$  or  $n = -1$ ; it is only defined for non-negative integers.
- The following is a rudimentary implementation of `fib(n)` in C:

```

/* fib: Returns the n-th Fibonacci term for non-negative n; returns -1 on
   bad input or overflow */
long fib(long n) {
    if (n <= 0)    # Bad input
        return -1;
    long f0 = 0;
    long f1 = 1;
    for (long m = 1; m < n && f1 >= f0; m++) {
        long temp = f1;
        f1 += f0;
        f0 = temp;
    }
}

```

```

    if (f1 <= f0)      # Overflow
        return -1;
    return f1;
}

```

---

(c)

```

.globl main
main:
    addi    a0,zero,9
    jal     fib
    jr      ra
    # long fib(long n);
fib:
    bge     zero,a0,.failed # if n <= 0 goto .failed
    addi    t0,zero,0      # f0 = 0
    addi    t1,zero,1      # f1 = 1
    addi    t2,zero,1      # m = 1
.loop:
    bge     t2,a0,.end     # if m >= n
    blt     t1,t0,.failed  # if f0 >= f1 goto .end
    addi    t3,t1,0        # temp = f1;
    add     t1,t1,t0        # t1 = t1 + t0
    addi    t0,t3,0        # t0 = temp
    addi    t2,t2,1        # m++
    j       .loop          # repeat
.end:
    bge     t0,t1,.failed  # if f0 >= f1 goto .failed
    addi    a0,t1,0        # set f1 as return value
    jr      ra             # return
.failed:
    addi    a0,zero,-1     # return -1;
    jr      ra

```

---

When run at <https://venus.kvakil.me/>, it placed 0x22 in register a0, which gives 34.

**Exercise 6.20.** Consider Code Example 6.28. For this exercise, assume *factorial*(*n*) is called with input argument *n* = 5.

- (a) What value is in a0 when `factorial` returns to the calling function?
- (b) Suppose you replace the instructions at address 0x8508 and 0x852C with nops. Will the program:
  - (1) Enter an infinite loop but not crash;
  - (2) crash (cause the stack to grow or shrink beyond the dynamic data segment or the PC to jump to a location outside the program);

- (3) produce an incorrect value in `a0` when the program returns to loop (if so, what value?); or
  - (4) run correctly despite the deleted lines?
- (c) Repeat part (b) with the following instruction modifications:
- (1) Replace the instructions at addresses `0x8504` and `0x8528` with nops.
  - (2) Replace the instruction at address `0x8518` with a nop.
  - (3) Replace the instruction at address `0x8530` with a nop.

**Solution:** The code for factorial is below:

---

```
factorial:
0x8500: addi    sp, sp,    -8      # make room for a0, ra
0x8504: sw      a0, 4(sp)
0x8508: sw      ra, 0(sp)
0x850C: addi    t0, zero,  1      # temporary = 1
0x8510: bgt     a0, t0,    else   # if n>1 goto else
0x8514: addi    a0, zero,  1      # otherwise return 1
0x8518: addi    sp, sp,    8      # restore sp
0x851C: jr      ra              # return
else:
0x8520: addi    a0, a0,    -1     # n = n - 1
0x8524: jal     factorial        # recursive call
0x8528: lw      t1, 4(sp)        # restore n into t1
0x852C: lw      ra, 0(sp)        # restore ra
0x8530: addi    sp, sp,    8      # restore sp
0x8534: mul     a0, t1,    a0     # a0 = n * factorial(n-1)
0x8538: jr      ra              # return
```

---

- (a) `a0` will have the value  $5! = 120$ .
- (b) Instruction `0x8508` pushes the `ra` register onto the stack in case the function recursively calls itself. Instruction `0x8528` restores loads `n` from the stack onto register `t1`.

When the function is called with `n` set to 5, it calls itself recursively without storing `ra` on the stack. However, the `jal factorial` instruction at each recursion level overwrites `ra`. This means that when it reaches the base case, calling itself with `n` set to 1, it returns to the level where `n` is 2, and then the instruction at address `0x852C`, namely `lw ra, 0(sp)`, will fail to find an address at that location. Since that stack address was not initialized (as a consequence of making `0x8058` a nop), it may have an existing value that is invalid, causing the program to crash because the PC will jump to a location outside the program.

- (c) If `0x8504` and `0x8528` are replaced with nops, then when we get to the instruction at address `0x8534` to compute `n * factorial(n-1)`, register `t1` will not have the expected value `n`. Suppose that `t1` had value  $x$  before `factorial(5)` was called. An incorrect



value will be produced in **a0** when all **factorial(5)** returns, namely  $x^4$ , assuming no overflow.

If the instruction at **0x8518** is replaced with a **nop**, then the base case will not restore the stack pointer. When the instruction at address **0x852C** is executed, **factorial(2)** will return to address **0x8528** again as it should. However, **factorial(2)** will return 1 instead of 2, and similarly, **factorial(3)** will return 2, **factorial(4)** will return 6, and instead of ending, **factorial(5)** will jump again to **0x8528**, computing one more iteration. This time it will find 5 at **4(sp)**. Surprisingly, it will compute the correct return value.

The instruction at **0x8530** restores the stack pointer when a non-leaf instance of **factorial** is about to return. This means that **factorial(2)** will return to **factorial(3)** but **factorial(3)** will see the same stack values that **factorial(2)** saw. The code results in an infinite loop.

**Exercise 6.22.** Convert the following RISC-V assembly code into machine language. Write the instruction in hexadecimal.

---

```
addi    s3, s4, 28
sll     t1, t2, t3
srli    s3, s1, 14
sw      s9, 16(t4)
```

---

**Solution:** The **addi** instructions is of *I*-type, which has the format

I-Type				
imm <sub>11:0</sub>	rs1	funct3	rd	op
12 bits	5 bits	3 bits	5 bits	7 bits

The opcode and function field for **addi** are **op** = 19(0010011<sub>2</sub>), and **funct<sub>3</sub>** = 0(000<sub>2</sub>). The source register is **s4** = **x20**, so **rs1** = 20(10100<sub>2</sub>), and the destination register is **s3** so **rd** = 19(10011<sub>2</sub>). Finally, we can represent 28 as a 12-bit immediate with the encoding 000000011100<sub>2</sub>. Altogether the encoding is

$$0000\ 0001\ 1100\ 1010\ 0000\ 1001\ 1001\ 0011_2 = 0x01CA0993$$

Next, **sll** is of *R*-type, which has the format:

R-type					
funct7	rs2	rs1	funct3	rd	op
7 bits	5 bits	5 bits	3 bits	5 bits	7 bits

For **sll** we have **opcode** = 0110011<sub>2</sub>, **funct3** = 001<sub>2</sub>, and **funct7** = 0000000<sub>2</sub>. Here, **t2** = **x7** and **t3** = **x28** are source registers **rs1** and **rs2**, respectively, and **t3** = **x6**. So **rs1** = 00111<sub>2</sub>, **rs2** = 11100<sub>2</sub>, and **rd** = 00110<sub>2</sub>. The full encoding is:

$$0000\ 0001\ 1100\ 0011\ 1001\ 0011\ 0011\ 0011_2 = 0x01C39333$$

Next, `srli` is of *I*-type. It has **opcode** = 0010011<sub>2</sub> and **funct3** = 101<sub>2</sub>. The destination is `s3 = x19` and the source is `s1 = x9`, so **rd** = 10011<sub>2</sub>, and **rs1** = 10001<sub>2</sub>. The immediate is 14 = 000000001110<sub>2</sub>. The encoding is

$$0000\ 0000\ 1110\ 1000\ 1101\ 1001\ 1001\ 0011_2 = 0x00E8D993$$

Finally we consider `sw` of *S*-type, which has format:

S-type					
imm <sub>11:5</sub>	rs2	rs1	funct3	imm <sub>4:0</sub>	op
7 bits	5 bits	5 bits	3 bits	5 bits	7 bits

`sw` has **opcode** = 0100011<sub>2</sub> and **funct3** = 010<sub>2</sub>. Moreover, `s9 = x25` and `t4 = x29` are source registers `rs2` and `rs1`, respectively. Therefore **rs2** = 11001<sub>2</sub> and **rs1** = 11101<sub>2</sub>. The immediate is 16, so we have **imm**<sub>4:0</sub> = 10000<sub>2</sub> and **imm**<sub>11:5</sub> = 0000000<sub>2</sub>. The encoding is:

$$0000\ 0001\ 1001\ 1110\ 1010\ 1000\ 0010\ 0011_2 = 0x019EA823$$

The translation for all four instructions is:

---

```
0x01CA0993 # addi s3, s4, 28
0x01C39333 # sll t1, t2,t3
0x00E8D993 # srli s3, s1, 14
0x019EA823 # sw s9, 16(t4)
```

---

**Exercise 6.23.** Repeat Exercise 6.22 for the following RISC-V assembly code:

---

```
add    s7, s8, s9
srai   t0, t1, 0xC
ori     s3, s1, 0xABC
lw      s4, 0x5C(t3)
```

---

**Solution:** `add` is of *R*-type. It has **opcode** = 0110011<sub>2</sub>, **funct7** = 0000000<sub>2</sub>, **funct3** = 000<sub>2</sub>. The instruction has source registers `s8 = x24` and `s9 = x25`, and destination register `s7 = x23`. Therefore we have **rs1** = 11000<sub>2</sub>, **rs2** = 11001<sub>2</sub>, and **rd** = 10111<sub>2</sub>. The encoding is

$$0000\ 0001\ 1001\ 1100\ 0000\ 1011\ 1011\ 0011_2 = 0x019C0BB3$$

Next, `srai` is of *I*-type. It has **opcode** = 0010011<sub>2</sub> and **funct3** = 101<sub>2</sub>. The instruction has destination `t0 = x5` and source `t1 = x6`, so **rd** = 00101<sub>2</sub> and **rs1** = 00110<sub>2</sub>. Finally the immediate is 0xC, so **imm**<sub>11:0</sub> = 000000001100<sub>2</sub>. The full instruction encoding is

$$0000\ 0000\ 1100\ 0011\ 0101\ 0010\ 1001\ 0011_2 = 0x00C35293$$

Next, `ori` is also of *I*-type. It has **opcode** = 0010011<sub>2</sub> and **funct3** = 110<sub>2</sub>. The instruction has destination `s3 = x19` and source `s1 = x9`, so **rd** = 10011<sub>2</sub> and **rs1** = 10001<sub>2</sub>. Finally, the immediate is 0xABC, so **imm**<sub>11:0</sub> = 1010 1011 1100<sub>2</sub>. The encoding is

$$1010\ 1011\ 1100\ 1000\ 1110\ 1001\ 1001\ 0011_2 = 0xABC8E993$$

Finally, `lw` is of *I*-type. It has **opcode** = 0000011<sub>2</sub> and **funct3** = 010<sub>2</sub>. The instruction has destination **s4** = **x20** and source **t3** = **x28**, so **rd** = 10100<sub>2</sub> and **rs1** = 11100<sub>2</sub>. Finally, the immediate is 0x5C, so **imm**<sub>11:0</sub> = 0000 0100 1100<sub>2</sub>. The encoding is

$$0000\ 0100\ 1100\ 1110\ 0010\ 1010\ 0000\ 0011_2 = 0x04CE2A03$$

The translation for all four instructions is:

---

```
0x019C0BB3 # add  s7, s8, s9
0x00C35293 # srai t0, t1, 0xC
0xABC8E993 # ori  s3, s1, 0xABC
0x04CE2A03 # lw  s4, 0x5C(t3)
```

---

**Exercise 6.24.** Consider all instructions with an immediate field.

- Which instructions from Exercise 6.22 use an immediate field in their machine code format?
- What is the instruction type (*I*-, *S*-, *B*-, *U*-, or *J*-type) for the instructions from part (a)?
- Write out the 5- to 21-bit immediates of each instruction from part (a) in hexadecimal. If the number is extended, also write them as 32-bit extended immediates. Otherwise, indicate that they are not extended.

**Solution:**

- The following instructions in Exercise 6.22 use an immediate field in their machine code format:

---

```
addi    s3, s4, 28
srli    s3, s1, 14
sw      s9, 16(t4)
```

---

- The `addi` and `srli` instruction are of *I*-type, and the `sw` instruction is of *S*-type.
- 

---

```
0x01C   # addi s3, s4, 28
0x00E   # srli s3, s1, 14
0x012   # sw   s9, 16(t4)
```

---

**Exercise 6.25.** Repeat Exercise 6.24 for the instructions in Exercise 6.23.

**Solution:**

- The following instructions in Exercise 6.23 use an immediate field in their machine code format:

---

```

srai    t0, t1, 0xC
ori     s3, s1, 0xABC
lw      s4, 0x5C(t3)

```

---

(b) All three instructions are of *I*-type.

(c)

---

```

0x00C      # srai t0, t1, 0xC
0xFFFFFABC # ori  s3, s1, 0xABC
0x04C      # lw   s4, 0x5C(t3)

```

---

**Exercise 6.26.** Consider the RISC-V machine code snippet below. The first instruction is listed at the top.

---

```

0x01800513
0x00300593
0x00000393
0x00058E33
0x01C54863
0x00138393
0x00BE0E33
0xFF5FF06F
0x00038533

```

---

- (a) Convert the machine code snippet into RISC-V assembly language.
- (b) Reverse engineer a high-level program that would compile into this assembly language routine and write it. Clearly comment your code.
- (c) Explain in words what the program does. **a0** and **a1** are the inputs, and they initially contain positive numbers *A* and *B*. At the end of the program, register **a0** holds the output (i.e., return value).

**Exercise 6.27.** Repeat Exercise 6.26 for the following machine code. **a0** and **a1** are the inputs. **a0** contains a 32-bit number and **a1** is the address of a 32-element array of characters (**char**).

---

```

0x01F00393
0x00755E33
0x001E7E13
0x01C580A3
0x00158593
0xFFF38393
0xFE03D6E3
0x00008067

```

---

**Exercise 6.28.** Convert the following branch instructions into machine code. Instructions addresses are given to the left of each instruction.

(a)

---

0x0000A000	beq t4, zero, Loop
0x0000A004	...
0x0000A008	...
0x0000A00C Loop:	...

---

(b)

---

0x00801000	bne s5, a1, L1
...	...
0x0080174C L1:	...

---

(c)

---

0x0000C10C Back:	...
...	...
0x0000D000	blt s1, s2, Back

---

**Exercise 6.29.** Convert the following branch instructions into machine code. Instruction addresses are given to the left of each instruction.

(a)

---

0xAA00E124	blt t4, s3, Loop
0xAA00E128	...
0xAA00E12C	...
0xAA00E130 Loop:	...

---

(b)

---

0xC0901000	bge t1, t2, L1
...	...
0xC090174C	...

---

(c)

---

0x1230D10C Back:	...
...	...
0x1230D908	bne s10, s11, Back

---

(d)

0xAB0C99A8	beq s0, s1, L2
...	...
0xAB0CA0FC L2:	...

(e)

0xFFABCF04 L3:	...
...	...
0xFFABD640	blt s1, t3, L3

**Exercise 6.30.** Convert the following jump instructions into machine code. Instruction addresses are given to the left of each instruction.

(a)

0x123ABC0	j Loop
...	...
0x123CABBC Loop:	...

(b)

0x12345678 Back	...
...	...
0x123B8760	jal s0, Back

(c)

0xAABBCCD0	jal L1
...	...
0xAABDCD98 L1:	...

(d)

0x11223344	j L2
...	...
0x1127BCDC L2:	...

(e)

0x9876543C L3:	...
...	...
0x9886543C	jal L3